

Sam Kritchevsky

Location: New York, NY

[GitHub](#) | [LinkedIn](#) | sam.kritch@gmail.com | 336-414-8694

ABOUT ME

I'm a backend engineer and "data generalist", with 5 years' experience working across a modern cloud-based data stack in Python, Scala, and SQL, and a strong math background. I'm returning to work after a personal sabbatical, looking for a role at an energetic small-/medium-sized company where I can quickly be productive and strategically impactful, and in opportunities to develop skills in adjacent technical areas such as ML or full-stack engineering.

TECHNICAL SKILLS

Languages	: Python, Scala, SQL, some Javascript / Typescript
Technologies	: PostgreSQL, Redshift, Elasticsearch, Redis, PostGIS, Spark, Flink, RabbitMQ, Akka, DBT, Docker, Nomad, Datadog, Grafana, AWS, EMR, Lambda, Kinesis, S3, Neptune, Jupyter, Luigi, Dagster
Roles	: Data Engineering, Data Science, Data Platform, Backend Engineering, Analytics, ML Engineering

EXPERIENCE

Seatgeek	2016-2021
<i>Sr. Software Engineer, Data Platform Tech Lead</i>	1.5 Years

Tech lead of team of six engineers. Projects included:

- Rewrite of our in-house event ingestion service in Flink + Kafka, adding features and scalability to handle bursty traffic. I advocated for the project, designed the system, and implemented much of the business logic. Functionality included:
 - * Fanouts to front page recommendations, push notifications, and Redshift, with plans to support an online feature store.
 - * Realtime sessionization, with handling of late-arriving data.
 - * Validation against a repository of JSONSchemas, with alerting in Datadog
- Rollout of a standardized Jupyter setup based on Docker, Papermill, and NBViewer.
- Deployment of Amundsen, an OSS data-discovery platform, with a custom AWS Neptune backend.
- Owned and maintained the core data stack of *Druzhba* (a now-outdated Python ETL tool; I was the primary maintainer and open-sourced the project), DBT, Looker, Luigi, Redshift, and other AWS products.

<i>Sr. Data Scientist, Data Scientist II</i>	2.5 Years
<i>Data Scientist II</i>	

Specialized in user recommendation signals and our event catalog, working in Spark, SQL, and Python. Projects included:

- Algorithms for the weekly newsletter, cart abandonment and price drops notifications, in Spark + SQL. (This work tripled the revenue associated with the weekly newsletter.)
- Event and performer popularity models for event recommendations and search, built with Keras
- Entity-linking algorithm deduplicating users and device for marketing attribution and funnel KPIs.
- Ongoing involvement in design and measurement of metrics for recommendations, search, and CRM marketing.

<i>Software Engineer, Discovery</i>	1 Year
Contributed features for search, CRM marketing, and frontpage recommendations.	

Other Experience

<i>Greater Harlem Coalition (occasional volunteer)</i>	2022
Contributed data analysis, writing, editing, and website administration for a community advocacy group. My data pipeline uses DBT, Python, Dagster, and PostGIS (code here .)	
<i>Physics Teaching Assistant</i>	2014-2016

EDUCATION

University of Wisconsin	Madison, WI
<i>Coursework towards Physics Ph.D.</i>	2013-2015
University of North Carolina	Chapel Hill, NC
<i>B.S. Physics</i>	2012